

## June 2022

Search posts from June 2022

Search

17 posts: [8 entries](#), [7 links](#), [2 quotes](#)

### June 5, 2022

## How to use the GPT-3 language model

**Playground**

Load a preset...

Save

View code

Share

...

Three reasons to start a succulent garden

1. Succulents are low-maintenance: They don't require much watering or fertilizing, and they can tolerate a wide range of light conditions.
2. Succulents are drought-tolerant: They're perfect for areas that receive little rainfall or irrigation.
3. Succulents add interest and variety to the landscape: With their wide range of shapes, sizes, and colors, they can provide a unique and eye-catching addition to any garden.

Mode

Engine

text-davinci-002

Temperature

0.7

Maximum length

256

Stop sequences

Enter sequence and press Tab

Top P

1

Submit

112

Show All

I ran [a Twitter poll](#) the other day asking if people had tried [GPT-3](#) and why or why not. The winning option, by quite a long way, was “No, I don’t know how to”. So here’s how to try it out, for free, without needing to write any code.

[... [838 words](#)]

---

[5:28 pm](#) / [machine-learning](#), [ai](#), [gpt-3](#), [openai](#), [prompt-engineering](#), [generative-ai](#), [llms](#)

---

### June 7, 2022

## Weeknotes: Datasette Cloud ready to preview

I made an absolute ton of progress building Datasette Cloud on Fly this week, and also had a bunch of fun playing with GPT-3.

[... [370 words](#)]

---

[1 am](#) / [projects](#), [datasette](#), [weeknotes](#), [datasette-cloud](#), [fly](#), [gpt-3](#), [llms](#)

---

**[June 8, 2022](#)**

**[Announcing Pyston-lite: our Python JIT as an extension module](#)** ([via](#)) The Pyston JIT can now be installed in any Python 3.8 virtual environment by running “pip install pyston\_lite\_autoload”—which includes a hook to automatically inject the JIT. I just tried a very rough benchmark against Datasette (ab -n 1000 -c 10) and got 391.20 requests/second without the JIT compared to 404.10 request/second with it.

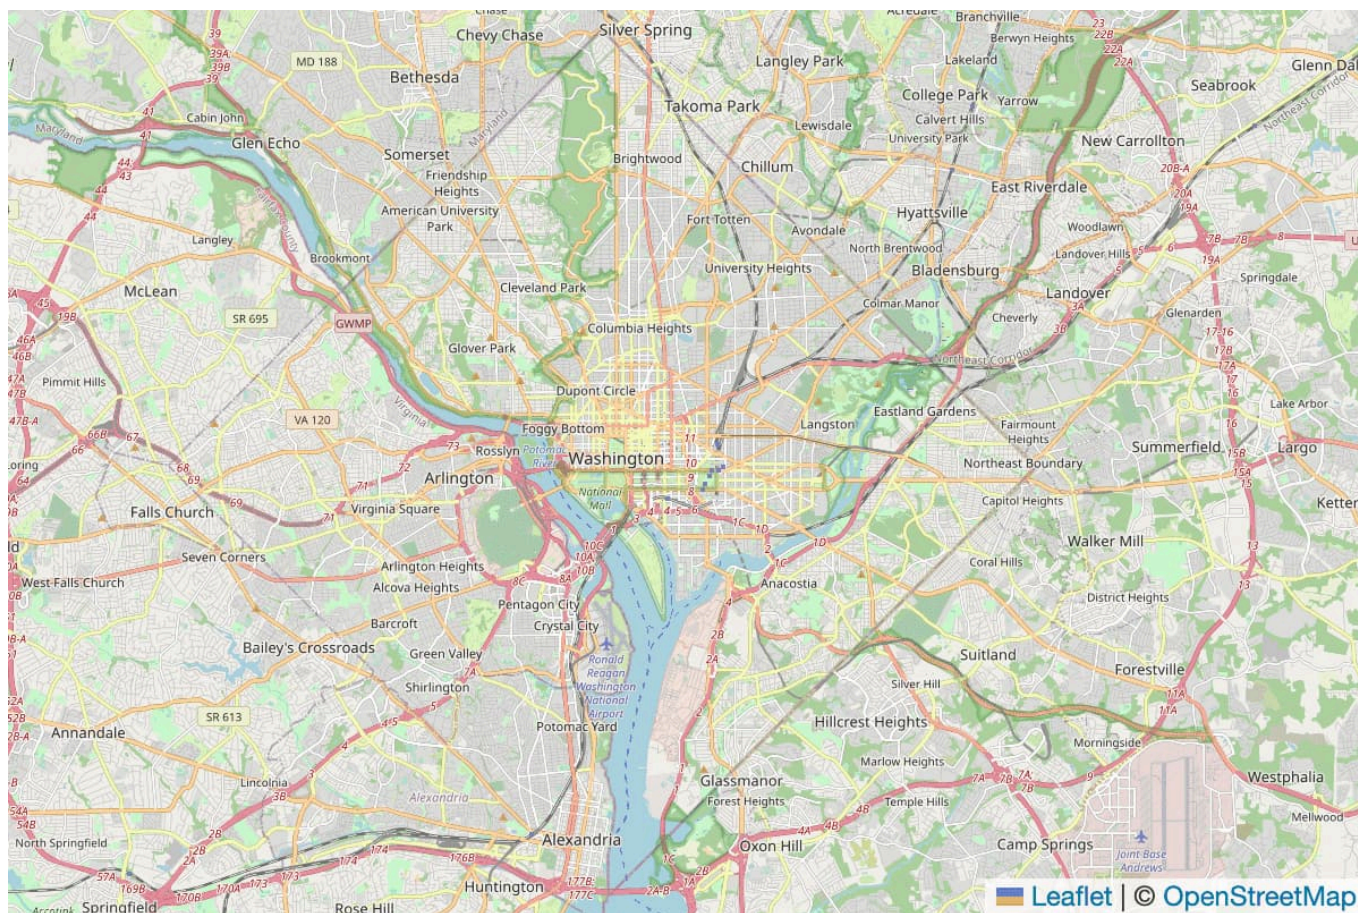
[# 5:58 pm](#) / [jit](#), [performance](#), [python](#)

**[The End of Localhost](#)**. swyx makes the argument for cloud-based development environments, and points out that many large companies—including Google, Facebook, Shopify and GitHub—have made the move already. I was responsible for the team maintaining the local development environment experience at Eventbrite for a while, and my conclusion is that with a large enough engineering team someone will ALWAYS find a new way to break their local environment: the idea of being able to bootstrap a fresh, guaranteed-to-work environment in the cloud at the click of a button could save SO much time and money.

[# 6:09 pm](#) / [software-engineering](#), [swyx](#)

**[June 12, 2022](#)**

## **[A tiny web app to create images from OpenStreetMap maps](#)**



Earlier today I found myself wanting to programmatically generate some images of maps.

[... [1,388 words](#)]

[5:49 am](#) / [gis](#), [maps](#), [openstreetmap](#), [projects](#), [shot-scraper](#)



# Simon Willison's Weblog

PHP, Python, CSS, XML and general web development

## This Site

[Front Page](#)  
[Archive](#)  
[Blogmarks](#)  
[Categories](#)  
[Comments](#)  
[Contact Me](#)

## Feeds

[Blog entries RSS](#)  
[Blog comments RSS](#)  
[Blogmarks RSS](#)

## Search Site

## Blogs I read

[Service Unavailable](#)

## Of Interest

[Incutio](#)  
[css-discuss wiki](#)  
[css-discuss archive](#)  
[XFML.org](#)  
[blo.gs](#)  
[Firebird](#)

## Tamarin

[3 days, 18 hours ago]

On Tuesday, the Mozilla Foundation and Adobe [announced](#) the [Tamarin project](#), an open-source ECMAScript virtual machine based on the ActionScript engine used by Flash Player 9.

Frank Hecker's [overview](#) of what this means is useful, but [the Tamarin source code itself](#) provides this interesting piece of historical insight:

AVM+ is the ActionScript Virtual Machine

AVM+ offers an order of magnitude performance increase over the "Classic AVM" in Flash Player 7. Our performance target is 10X.

AVM+ implements ActionScript 3.0, the new version of the ActionScript language that is compliant with the ECMAScript Edition 4 standard.

AVM+ is also built for modularity. It will be part of the Flash Player, but is a self-contained module which can be incorporated into other programs with ease. It may also be submitted to the ECMA standards organization as a reference implementation of ECMAScript Edition 4.

## Blogmarks:

- Concerned is over
- A Second Life teaser
- Dunstan Orchard working for Flickr
- Y! Cool Thing: Mapping The Globe
- Yahoo! bookmarks uses symfony
- How the news breaks
- Tamarin Project (via)
- Stackless Python Continues to Amaze
- PHP: JSON Functions
- Cleaner Callbacks With Partial Application
- Python Standard Encodings
- The Django Book
- A9 Maps are no longer available
- Real-Time 3D in Javascript
- Rafe Colburn: The secret of YouTube
- .. and [2365 more](#) [\[RSS\]](#)

I started this blog on [June 12th 2002](#)—twenty years ago today! To celebrate two decades of blogging, I decided to pull together some highlights and dive down a self-indulgent nostalgia hole.

[... [4,455 words](#)]

---

[10:59 pm](#) / [adrian-holovaty](#), [blogging](#), [personalnews](#)

---

## [June 13, 2022](#)

[Making Code Faster](#). Tim Bray's detailed guide to using the Go profiler.

# [7:40 pm](#) / [go](#), [profiling](#), [tim-bray](#)

---

## [June 18, 2022](#)

Becoming a good engineer is about collecting experience. Each project, even small ones, is a chance to add new techniques and tools to your toolbox. Where this delivers even more value is when you can solve problems by pairing techniques learned on one project with tools learned working on another. It all adds up.

— [Addy Osmani](#)

## **June 19, 2022**

### **Weeknotes: datasette-socrata, and the last 10%...**

... takes 90% of the work. I continue to work towards a preview of the new Datasette Cloud, and keep finding new “just one more things” to delay inviting in users.

[... [1,214 words](#)]

---

[3:26 am](#) / [open-data](#), [plugins](#), [datasette](#), [weeknotes](#), [datasette-cloud](#), [sqlite-utils](#), [annotated-release-notes](#)

---

**WarcDB** ([via](#)) Florents Tselai built this tool for loading web crawl data stored in WARC (Web ARChive) format into a SQLite database for smaller-scale analysis with SQL, on top of my [sqlite-utils](#) Python library.

# 6:08 pm / [archiving](#), [sqlite](#), [sqlite-utils](#)

---

## **June 20, 2022**

**The State of WebAssembly 2022**. Colin Eberhardt talks through the results of the State of WebAssembly 2022 survey. Rust continues to dominate as the most popular language for working to WebAssembly, but Python has a notable increase of interest.

# 6:07 pm / [rust](#), [webassembly](#)

---

## **Joining CSV files in your browser using Datasette Lite**

# data

Custom SQL query returning more than 1,000 rows ([hide](#))

```
select
  [us-counties-recent].*,
  us_census_county_populations_2019.population,
  1.0 * [us-counties-recent].cases /
us_census_county_populations_2019.population * 1000000 as cases_per_million
from
  [us-counties-recent]
  join us_census_county_populations_2019 on
us_census_county_populations_2019.fips = [us-counties-recent].fips
where
  date = (select max(date) from [us-counties-recent])
order by
  cases_per_million desc
```

Run SQL

This data as [json](#), [CSV](#)

date	county	state	fips	cases	deaths	population	cases_per_million
2022-06-19	Loving	Texas	48301	218	1	169	1289940.8284023667
2022-06-19	Chattahoochee	Georgia	13053	7515	22	10907	689007.0596864399
2022-06-19	Northwest Arctic Borough	Alaska	2188	4927	14	7621	646503.0835848314
2022-06-19	Nome Census Area	Alaska	2180	6306	5	10004	630347.8608556577
2022-06-19	Bethel Census Area	Alaska	2050	10670	41	18386	580332.8619601871
2022-06-19	Dimmit	Texas	48127	5828	51	10124	575661.7937574082
2022-06-19	Crowley	Colorado	8025	3378	31	6061	557333.7733047352
2022-06-19	Jim Hogg	Texas	48247	2786	22	5200	535769.2307692308

I added a new feature to [Datasette Lite](#)—my version of [Datasette](#) that runs entirely in your browser using WebAssembly ([previously](#)): you can now use it to load one or more CSV files by URL, and then run SQL queries against them—including joins across data from multiple files.

[... [546 words](#)]

9:20 pm / [csv](#), [projects](#), [sql](#), [datasette](#), [webassembly](#), [datasette-lite](#), [cors](#)

## June 23, 2022

[How Imagen Actually Works](#). Imagen is Google’s new text-to-image model, similar to (but possibly even more effective than) DALL-E. This article is the clearest explanation I’ve seen of how Imagen works: it uses Google’s existing T5 text encoder to convert the input sentence into an encoding that captures the semantic meaning of the sentence (including things like items being described as being on top of other items), then uses a trained diffusion model to generate a 64x64 image. That image is passed through two super-res models to increase the resolution to the final 1024x1024 output.

# 6:05 pm / [google](#), [machine-learning](#), [ai](#), [dalle](#), [generative-ai](#)

## First impressions of DALL-E, generating images from text





I made it off the [DALL-E waiting list](#) a few days ago and I've been having an enormous amount of fun experimenting with it. Here are some notes on what I've learned so far (and a bunch of example images too).

[... [2,102 words](#)]

---

11:05 pm / [machine-learning](#), [ai](#), [openai](#), [dalle](#), [prompt-engineering](#), [generative-ai](#), [text-to-image](#)

---

**[June 28, 2022](#)**

**[The Magic Interview Question](#)** ([via](#)) Jeff Gothelf explains why “Tell me about the last time you [did something]” is the most valuable question you can ask when interviewing a user or potential user.

# [2:26 pm](#) / [usability](#)

---

The general idea of an “Islands” architecture is deceptively simple: render HTML pages on the server, and inject placeholders or slots around highly dynamic regions. These placeholders/slots contain the server-rendered HTML output from their corresponding widget. They denote regions that can then be "hydrated" on the client into small self-contained widgets, reusing their server-rendered initial HTML.

— [Jason Miller](#)

# [3:01 pm](#) / [javascript](#)

[June 30, 2022](#)

## [s3-ocr: Extract text from PDF files stored in an S3 bucket](#)

pages

Data source: [Library of Congress](#) · About: [s3-ocr demo](#)

147 rows where search matches "houdini"

Search:

houdini

- column -

=

Apply

[View and edit SQL](#)

This data as [json](#), [CSV \(advanced\)](#)

path 2 ✖

[unmaskingrobert00houdgoog.pdf](#)

146

[latestmagicbeing00hoff.pdf](#)

1

Link

path ⚙

page ⚙

folder ⚙

text ⚙

[latestmagicbeing00hoff.pdf,238](#)

latestmagicbeing00hoff.pdf

238

L'ENVOI  
WITH these last lines I lay down my pen, as I  
have long since laid down the wand. I do SO with  
regret, for writing about magic has always been to

I've released [s3-ocr](#), a new tool that runs Amazon's [Textract](#) OCR text extraction against PDF files in an S3 bucket, then writes the resulting text out to a SQLite database with full-text search configured so you can run searches against the extracted data.

[... [1,493 words](#)]

[9:40 pm](#) / [aws](#), [ocr](#), [pdf](#), [projects](#), [s3](#), [weeknotes](#), [s3-credentials](#)

[2022](#) » June

M	T	W	T	F	S	S
		1	2	3	4	5
6	7	8	9	10	11	12
13	14	15	16	17	18	19
20	21	22	23	24	25	26
27	28	29	30			

