

November 2022

29 posts: [8 entries](#), [18 links](#), [3 quotes](#)

Nov. 1, 2022

[RFC 7807: Problem Details for HTTP APIs](#) ([via](#)) This RFC has been brewing for quite a while, and is currently in last call (ends 2022-11-03). I'm designing the JSON error messages for Datasette at the moment so this could not be more relevant for me.

[3:15 am](#) / [errors](#), [http](#), [json](#), [mark-nottingham](#), [rfc](#), [standards](#)

Nov. 4, 2022

[Don't Read Off The Screen](#) ([via](#)) Stuart Langridge provides a fantastic set of public speaking tips in a five minute lightning talk remix of Sunscreen. Watch with sound.

[4:02 pm](#) / [speaking](#), [stuart-langridge](#)

Nov. 5, 2022

[Nikodemus' Guide to Mastodon](#) ([via](#)) I've been reading a bunch of different Mastodon guides and this one had pretty much exactly the information I needed to see when I first started out.

[4:18 am](#) / [mastodon](#)

It looks like I'm moving to Mastodon

Elon Musk laid off about half of Twitter this morning. There are [many terrible stories](#) emerging about how this went down, but one that particularly struck me was that he laid off [the entire accessibility team](#). For me this feels like a microcosm of the whole situation. Twitter's priorities are no longer even remotely aligned with my own.

[... [1,546 words](#)]

[5:32 am](#) / [accessibility](#), [twitter](#), [mastodon](#), [fediverse](#)

[GOV.UK: Rules for getting production access](#) ([via](#)) Fascinating piece of internal documentation on GOV.UK describing their rules, procedures and granted permissions for their deployment and administrative ops roles.

[6:25 pm](#) / [security](#), [gov-uk](#)

Nov. 6, 2022

What to blog about

You should start a blog. Having your own little corner of the internet is good for the soul!

[... [520 words](#)]

[5:05 pm](#) / [blogging](#), [writing](#)

Nov. 7, 2022

[Blessed.rs Crate List](#) ([via](#)) Rust doesn't have a very large standard library, so part of learning Rust is figuring out which of the third-party crates are the best for tackling common problems. This here is an opinionated guide to crates, which looks like it could be really useful.

[7:25 pm](#) / [rust](#)

Nov. 8, 2022

Mastodon is just blogs

And that's *great*. It's also the return of Google Reader!

[... [1,560 words](#)]

[3:48 pm](#) / [blogging](#), [google-reader](#), [twitter](#), [mastodon](#), [activitypub](#), [fediverse](#), [sidekiq](#)

Nov. 9, 2022

Designing a write API for Datasette

Building out [Datasette Cloud](#) has made one thing clear to me: Datasette needs a write API for ingesting new data into its attached SQLite databases.

[... [1,493 words](#)]

[7:44 pm](#) / [api-design](#), [projects](#), [datasette](#), [weeknotes](#), [datasette-cloud](#), [sqlite-utils](#)

[Inside the mind of a frontend developer: Hero section](#). Ahmad Shadeed provides a fascinating, hyper-detailed breakdown of his approach to implementing a "hero section" component using HTML and CSS, including notes on CSS grids and gradient backgrounds.

[7:54 pm](#) / [css](#), [ahmad-shadeed](#)

[Semantic text search using embeddings](#). Example Python notebook from OpenAI demonstrating how to build a search engine using embeddings rather than straight up token matching. This is a fascinating way of implementing search, providing results that match the intent of the search ("delicious beans" for example) even if none of the keywords are actually present in the text.

7:57 pm / [machine-learning](#), [search](#), [openai](#), [embeddings](#)

[PyScript Updates: Bytecode Alliance, Pyodide, and MicroPython](#). Absolutely huge news about Python on the Web tucked into this announcement: Anaconda have managed to get a version of MicroPython compiled to WebAssembly running in the browser. Pyodide weighs in at around 6.5MB compressed, but the MicroPython build is just 303KB—the size of a large image. This makes Python in the web browser applicable to so many more potential areas.

10:26 pm / [python](#), [webassembly](#), [pyodide](#)

[Nov. 11, 2022](#)

[Home invasion: Mastodon's Eternal September begins](#). Hugh Rundle's thoughtful write-up of the impact of the massive influx of new users from Twitter on the existing Mastodon community. If you're new to Mastodon (like me) you should read this and think carefully about how best to respectfully integrate with your new online space.

12:47 am / [mastodon](#)

[Nov. 13, 2022](#)

[Datasette is 5 today: a call for birthday presents](#)



Five years ago today I published the first release of Datasette, in [Datasette: instantly create and publish an API for your SQLite databases](#).

[... [548 words](#)]

7:27 pm / [projects](#), [datasette](#)

[Nov. 16, 2022](#)

[JSON Changelog with SQLite](#) ([via](#)) One of my favourite database challenges is how to track changes to rows over time. This is a neat recipe from 2018 which uses SQLite triggers and the SQLite JSON functions to serialize older versions of the

rows and store them in TEXT columns.

[# 3:41 am](#) / [databases](#), [json](#), [sqlite](#)

[fasiha/yamanote](#) ([via](#)) Yamanote is “a guerrilla bookmarking server” by Ahmed Fasih—it works using a bookmarklet that grabs a full serialized copy of the page—the innerHTML of both the head and body element—and passes it to the server, which stores it in a SQLite database. The files are then served with a Content-Security-Policy: `default-src 'self'` header to prevent stored pages from fetching ANY external assets when they are viewed.

[# 3:48 am](#) / [bookmarks](#), [sqlite](#)

These kinds of biases aren’t so much a technical problem as a sociotechnical one; ML models try to approximate biases in their underlying datasets and, for some groups of people, some of these biases are offensive or harmful. That means in the coming years there will be endless political battles about what the ‘correct’ biases are for different models to display (or not display), and we can ultimately expect there to be as many approaches as there are distinct ideologies on the planet. I expect to move into a fractal ecosystem of models, and I expect model providers will ‘shapeshift’ a single model to display different biases depending on the market it is being deployed into. This will be extraordinarily messy.

— [Jack Clark](#)

[# 11:04 pm](#) / [machine-learning](#), [ai](#), [generative-ai](#), [jack-clark](#), [llms](#)

[Nov. 18, 2022](#)

[Datasette Lite: Loading JSON data](#) ([via](#)) I added a new feature to Datasette Lite: you can now pass it the URL to a JSON file (hosted on a CORS-compatible hosting provider such as GitHub or GitHub Gists) and it will load that file into a database table for you. It expects an array of objects, but if your file has an object as the root it will search through it looking for the first key that is an array of objects and load those instead.

[# 6:43 pm](#) / [json](#), [projects](#), [datasette](#), [datasette-lite](#), [cors](#)

[Nov. 19, 2022](#)

... it [ActivityPub] is crucially good enough. Perfect is the enemy of good, and in ActivityPub we have a protocol that has flaws but, crucially, that works, and has a standard we can all mostly agree on how to implement - and eventually, I hope, agree on how to improve.

— [Andrew Godwin](#)

[# 4:02 pm](#) / [andrew-godwin](#), [mastodon](#), [activitypub](#)


[Nov. 20, 2022](#)

[Every remaining website using the .museum TLD](#) ([via](#)) Jonty did a survey of every one of the 1,134 domains using the .museum TLD, which dates back to 2001 and is managed by The Museum Domain Management Association.

[# 12:53 am](#) / [domains](#), [museums](#)

Tracking Mastodon user numbers over time with a bucket of tricks

Total users over time across all tracked instances



[Mastodon](#) is definitely having a moment. User growth is skyrocketing as more and more people migrate over from Twitter.

[... [1,534 words](#)]

[7 am](#) / [github](#), [projects](#), [datasette](#), [observable](#), [github-actions](#), [git-scraping](#), [git-history](#), [s3-credentials](#), [datasette-lite](#), [mastodon](#), [cors](#)

Nov. 21, 2022

[Building a BFT JSON CRDT](#) ([via](#)) Jacky Zhao describes their project to build a CRDT library for JSON data in Rust, and includes a thorough explanation of what CRDTs are and how they work. “I write this blog post mostly as a note to my past self, distilling a lot of what I’ve learned since into a blog post I wish I had read before going in”—the best kind of blog post!

[# 7:56 pm](#) / [json](#), [rust](#), [crdt](#)

Nov. 23, 2022

Weeknotes: Implementing a write API, Mastodon distractions

Everything is so *distracting* at the moment. The ongoing Twitter catastrophe, the great migration (at least amongst most of the people I pay attention to) [to Mastodon](#), the FTX calamity. It’s been very hard to focus!

[... [916 words](#)]

[4:57 am](#) / [datasette](#), [weeknotes](#), [mastodon](#)

Nov. 24, 2022

Microsoft Flight Simulator: WebAssembly ([via](#)) This is such a smart application of WebAssembly: it can now be used to write extensions for Microsoft Flight Simulator, which means you can run code from untrusted sources safely in a sandbox. I'm really looking forward to more of this kind of usage—I love the idea of finally having a robust sandbox for running things like plugins.

[2:08 am](#) / [microsoft](#), [webassembly](#)

Nov. 26, 2022

[An Interactive Guide to Flexbox](#). Joshua Comeau built this fantastic guide to CSS flexbox layouts, with interactive examples of all of the properties. This is a really useful tour of the layout model.

[2:50 am](#) / [css](#)

Coping strategies for the serial project hoarder

Massively increase your productivity on personal projects with comprehensive documentation and automated tests

Simon Willison, DjangoCon US 2022

Talk notes are linked from github.com/simonw

I gave a talk at DjangoCon US 2022 in San Diego last month about productivity on personal projects, titled “Massively increase your productivity on personal projects with comprehensive documentation and automated tests”.

[... [3,865 words](#)]

[3:47 pm](#) / [djangocon](#), [documentation](#), [productivity](#), [my-talks](#), [testing](#), [annotated-talks](#), [github-issues](#)

Nov. 28, 2022

If posts in a social media app do not have URLs that can be linked to and viewed in an unauthenticated browser, or if there is no way to make a new post from a browser, then that program is not a part of the World Wide Web in any meaningful way.

Consign that app to oblivion.

— [JWZ](#)

[6:22 am](#) / [jwz](#), [social-media](#)

[Nov. 29, 2022](#)

[Stable Diffusion 2.0 and the Importance of Negative Prompts for Good Results](#). Stable Diffusion 2.0 is out, and it’s a very different model from 1.4/1.5. It’s trained using a new text encoder (OpenCLIP, in place of OpenAI’s CLIP) which means a lot of the old tricks—notably using “Greg Rutkowski” to get high quality fantasy art—no longer work. What DOES work, incredibly well, is negative prompting—saying things like “cyberpunk forest by Salvador Dali” but negative on “trees, green”. Max Woolf explores negative prompting in depth in this article, including how to combine it with textual inversion.

[1:22 am](#) / [max-woolf](#), [stable-diffusion](#), [prompt-engineering](#), [generative-ai](#), [text-to-image](#)

[Scaling Mastodon: The Compendium](#) (via) Hazel Weakly’s collection of notes on scaling Mastodon, covering PostgreSQL, Sidekiq, Redis, object storage and more.

[5:46 am](#) / [postgresql](#), [redis](#), [scaling](#), [mastodon](#), [sidekiq](#)

[2022](#) » November

M	T	W	T	F	S	S
	1	2	3	4	5	6
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28	29	30				